Project 4 Notebook: Chocolate Sales Dataset (2023–2024)

This notebook presents the first three stages of an end-to-end data science workflow based on the Kaggle *Chocolate Sales Dataset (2023–2024)*.

It includes:

1. **Data Acquisition & Preparation**
   Load and merge multiple CSV files, clean the data, and construct a unified dataset.

2. **Exploratory Data Analysis (EDA)**
   Analyze distributions, trends, and relationships, and apply unsupervised learning to uncover hidden patterns.

3. **Feature Engineering & Preprocessing**
   Create new features, handle missing values, encode categorical variables, and scale numerical features to prepare for modeling.

1. Data Acquisition & Preparation

In [1]:
```
# Import required libraries
import pandas as pd
import numpy as np
from pathlib import Path

DATA_DIR = Path("./data")
```

In [2]:
```
# Load all CSV files into separate DataFrames
sales = pd.read_csv(DATA_DIR / "sales.csv")
products = pd.read_csv(DATA_DIR / "products.csv")
stores = pd.read_csv(DATA_DIR / "stores.csv")
customers = pd.read_csv(DATA_DIR / "customers.csv")
calendar = pd.read_csv(DATA_DIR / "calendar.csv")

print("sales:", sales.shape)
print("products:", products.shape)
print("stores:", stores.shape)
print("customers:", customers.shape)
print("calendar:", calendar.shape)
```

```
sales: (1000000, 11)
products: (200, 6)
stores: (100, 5)
customers: (50000, 5)
calendar: (731, 6)
```

In [3]:
```
# Function to summarize each table
def summarize(df, name):
    print(f"\n{name}")
```

```
    print("-" * 40)
    print("Shape:", df.shape)
    print("Missing values:")
    print(df.isna().sum())
    display(df.head(3))

summarize(sales, "sales")
summarize(products, "products")
summarize(stores, "stores")
summarize(customers, "customers")
summarize(calendar, "calendar")
```

```
sales
----------------------------------------
Shape: (1000000, 11)
Missing values:
order_id       0
order_date     0
product_id     0
store_id       0
customer_id    0
quantity       0
unit_price     0
discount       0
revenue        0
cost           0
profit         0
dtype: int64
```

	order_id	order_date	product_id	store_id	customer_id	quantity	unit_price
0	ORD00000001	2023-01-07	P0080	S093	C040749	5	14.43
1	ORD00000002	2023-10-22	P0173	S065	C020161	3	12.01
2	ORD00000003	2023-05-07	P0115	S078	C048069	2	10.02

```
products
----------------------------------------
Shape: (200, 6)
Missing values:
product_id       0
product_name     0
brand            0
category         0
cocoa_percent    0
weight_g         0
dtype: int64
```

	product_id	product_name	brand	category	cocoa_percent	weight_g
0	P0001	White Chocolate 80%	Mars	Truffle	80	120
1	P0002	Dark Chocolate 70%	Cadbury	Praline	70	100
2	P0003	Truffle Chocolate 70%	Hershey	Praline	70	120

```
stores
----------------------------------------
Shape: (100, 5)
Missing values:
store_id      0
store_name    0
city          0
country       0
store_type    0
dtype: int64
```

	store_id	store_name	city	country	store_type
0	S001	Chocolate Store 1	New York	Canada	Retail
1	S002	Chocolate Store 2	Melbourne	Canada	Mall
2	S003	Chocolate Store 3	Berlin	France	Mall

```
customers
----------------------------------------
Shape: (50000, 5)
Missing values:
customer_id       0
age               0
gender            0
loyalty_member    0
join_date         0
dtype: int64
```

	customer_id	age	gender	loyalty_member	join_date
0	C000001	40	Male	1	2025-05-21
1	C000002	47	Male	0	2021-12-26
2	C000003	58	Female	1	2022-09-13

```
calendar
----------------------------------------
Shape: (731, 6)
Missing values:
date           0
year           0
month          0
day            0
week           0
day_of_week    0
dtype: int64
```

	date	year	month	day	week	day_of_week
0	2023-01-01	2023	1	1	52	6
1	2023-01-02	2023	1	2	1	0
2	2023-01-03	2023	1	3	1	1

In [4]:

```python
# Data cleaning

# 1. Standardize column names
def clean_columns(df):
    df.columns = df.columns.str.strip().str.lower().str.replace(" ", "_")
```

```
    return df

sales = clean_columns(sales)
products = clean_columns(products)
stores = clean_columns(stores)
customers = clean_columns(customers)
calendar = clean_columns(calendar)

# 2. Convert date columns to datetime format
sales["order_date"] = pd.to_datetime(sales["order_date"])
customers["join_date"] = pd.to_datetime(customers["join_date"])
calendar["date"] = pd.to_datetime(calendar["date"])

# 3. Check for duplicate rows
print("Duplicate rows in sales:", sales.duplicated().sum())
print("Duplicate rows in products:", products.duplicated().sum())
print("Duplicate rows in stores:", stores.duplicated().sum())
print("Duplicate rows in customers:", customers.duplicated().sum())
print("Duplicate rows in calendar:", calendar.duplicated().sum())
```

```
Duplicate rows in sales: 0
Duplicate rows in products: 0
Duplicate rows in stores: 0
Duplicate rows in customers: 0
Duplicate rows in calendar: 0
```

In [5]:

```
# 4. Remove rows with missing key identifiers
sales = sales.dropna(subset=["product_id", "store_id", "customer_id", "order_dat

print("Cleaning completed.")
```

```
Cleaning completed.
```

In [6]:

```
# Merge all tables into a single DataFrame
df = sales.copy()

df = df.merge(products, on="product_id", how="left")
df = df.merge(stores, on="store_id", how="left")
df = df.merge(customers, on="customer_id", how="left")
df = df.merge(calendar, left_on="order_date", right_on="date", how="left")

print("Final dataset shape:", df.shape)
display(df.head())
```

```
Final dataset shape: (1000000, 30)
```

	order_id	order_date	product_id	store_id	customer_id	quantity	unit_price
0	0RD00000001	2023-01-07	P0080	S093	C040749	5	14.43
1	0RD00000002	2023-10-22	P0173	S065	C020161	3	12.01
2	0RD00000003	2023-05-07	P0115	S078	C048069	2	10.02
3	0RD00000004	2024-06-23	P0186	S088	C047901	2	14.66
4	0RD00000005	2024-09-24	P0197	S054	C033950	1	12.34

5 rows × 30 columns

In [7]:
```
# Check the merged dataset

# Display data types and non-null counts
print("Dataset info:")
df.info()

# Check missing values after merging
print("\nTop missing values:")
display(df.isna().sum().sort_values(ascending=False).head(10))
```

```
Dataset info:
<class 'pandas.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 30 columns):
 #   Column          Non-Null Count    Dtype
---  ------          --------------    -----
 0   order_id        1000000 non-null  str
 1   order_date      1000000 non-null  datetime64[us]
 2   product_id      1000000 non-null  str
 3   store_id        1000000 non-null  str
 4   customer_id     1000000 non-null  str
 5   quantity        1000000 non-null  int64
 6   unit_price      1000000 non-null  float64
 7   discount        1000000 non-null  float64
 8   revenue         1000000 non-null  float64
 9   cost            1000000 non-null  float64
 10  profit          1000000 non-null  float64
 11  product_name    990236 non-null   str
 12  brand           990236 non-null   str
 13  category        990236 non-null   str
 14  cocoa_percent   990236 non-null   float64
 15  weight_g        990236 non-null   float64
 16  store_name      1000000 non-null  str
 17  city            1000000 non-null  str
 18  country         1000000 non-null  str
 19  store_type      1000000 non-null  str
 20  age             1000000 non-null  int64
 21  gender          1000000 non-null  str
 22  loyalty_member  1000000 non-null  int64
 23  join_date       1000000 non-null  datetime64[us]
 24  date            1000000 non-null  datetime64[us]
 25  year            1000000 non-null  int64
 26  month           1000000 non-null  int64
 27  day             1000000 non-null  int64
 28  week            1000000 non-null  int64
 29  day_of_week     1000000 non-null  int64
dtypes: datetime64[us](3), float64(7), int64(8), str(12)
memory usage: 323.5 MB

Top missing values:
weight_g          9764
cocoa_percent     9764
category          9764
brand             9764
product_name      9764
gender               0
store_type           0
age                  0
loyalty_member       0
city                 0
dtype: int64
```

Data Summary and Preparation for analysis

The dataset used in this project is the *Chocolate Sales Dataset (2023–2024)* obtained from Kaggle. It consists of five relational tables: `sales`, `products`, `stores`, `customers`, and `calendar`.

The `sales` table serves as the main fact table, containing transaction-level data with one million records. It includes key variables such as order date, product ID, store ID, customer ID, quantity, price, and revenue-related metrics. The remaining tables provide additional contextual information:

- `products` contains product attributes such as brand, category, and cocoa percentage,
- `stores` includes location and store-type information,
- `customers` provides demographic details such as age, gender, and loyalty status,
- `calendar` offers structured time-related features such as year, month, and day of the week.

To prepare the data for analysis, several preprocessing steps were performed. First, column names were standardized to ensure consistency across tables. Date columns were converted into proper datetime format to enable time-based analysis. The data was checked for duplicate records, and none were found, so no deduplication was required. Missing values were identified, particularly in product-related fields after merging, which were attributed to unmatched keys between tables. These records were retained for now, as they represent a small proportion of the dataset.

Next, all tables were merged into a single unified dataset using appropriate keys (`product_id`, `store_id`, `customer_id`, and date fields). This resulted in a comprehensive transaction-level dataset enriched with product, store, customer, and temporal information.

The final dataset is well-structured and suitable for downstream tasks, including exploratory data analysis, feature engineering, and predictive modeling.

2. Exploratory Data Analysis (EDA)

In [8]:
```
# initial investigation
df.describe()
```

Out[8]:

	order_date	quantity	unit_price	discount	re
count	1000000	1000000.000000	1000000.000000	1000000.000000	1000000.0
mean	2023-12-31 19:03:34.991999	2.999589	9.002449	0.056248	25.4
min	2023-01-01 00:00:00	1.000000	3.000000	0.000000	2.4
25%	2023-07-02 00:00:00	2.000000	6.000000	0.000000	12.1
50%	2024-01-01 00:00:00	3.000000	9.010000	0.000000	21.9
75%	2024-07-02 00:00:00	4.000000	12.000000	0.150000	35.8
max	2024-12-31 00:00:00	5.000000	15.000000	0.200000	75.0
std	NaN	1.413621	3.462902	0.076809	16.3

The dataset contains a large number of transactions with varying sales volumes and pricing levels, indicating potential heterogeneity across products and customers.

In [9]:

```
# Aggregate total revenue by date
daily_sales = df.groupby("order_date")["revenue"].sum()

# Plot daily revenue trend
daily_sales.plot(figsize=(12,5), title="Daily Revenue Trend")
```

Out[9]: <Axes: title={'center': 'Daily Revenue Trend'}, xlabel='order_date'>

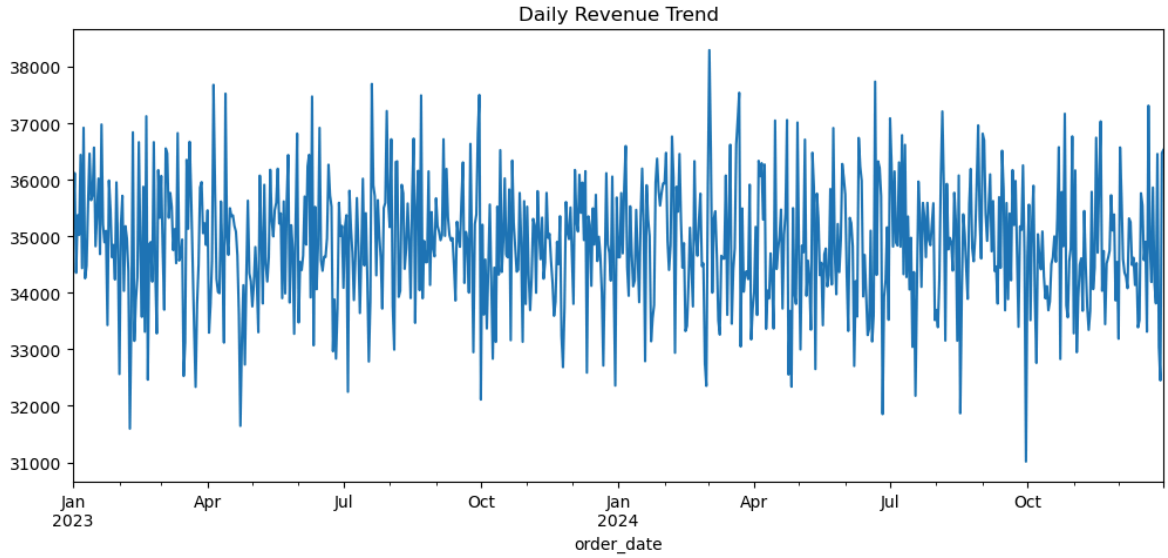

In [10]:

```
# Aggregate revenue by month
monthly_sales = df.groupby("month")["revenue"].sum()

# Plot monthly revenue
monthly_sales.plot(kind="bar", title="Revenue by Month")
```

Out[10]: <Axes: title={'center': 'Revenue by Month'}, xlabel='month'>

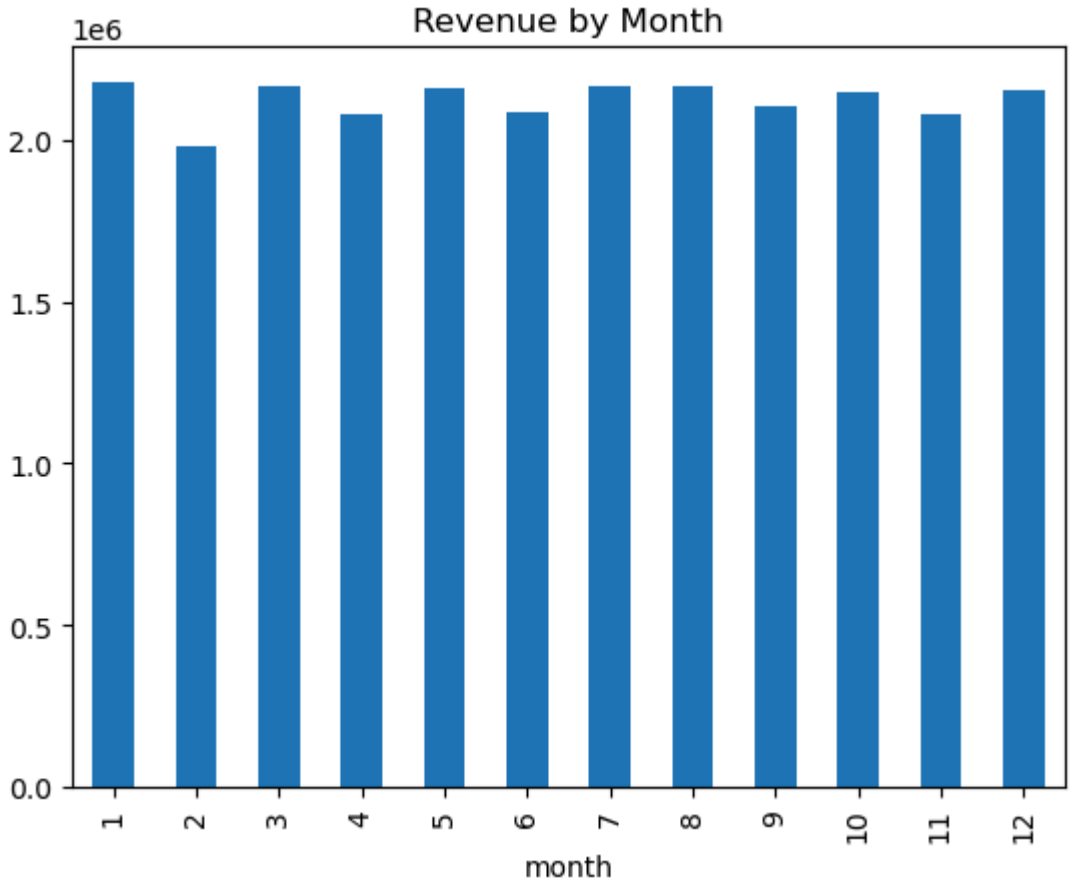

In [11]:
```
# Identify top-performing products by total revenue
top_products = df.groupby("product_name")["revenue"].sum().sort_values(ascending

# Plot top products
top_products.plot(kind="barh", title="Top 10 Products by Revenue")
```

Out[11]: <Axes: title={'center': 'Top 10 Products by Revenue'}, ylabel='product_name'>

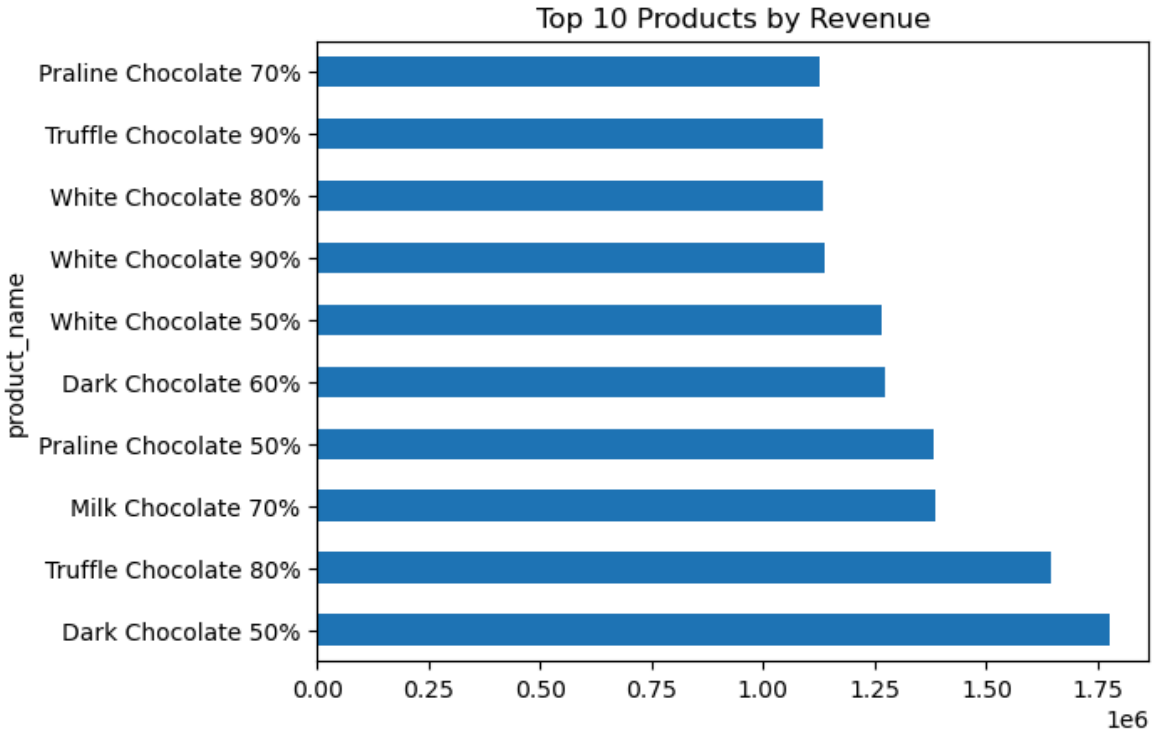

In [12]:
```
# Analyze average revenue by age group
df.groupby("age")["revenue"].mean().plot(title="Average Revenue by Age")
```

Out[12]:
```
<Axes: title={'center': 'Average Revenue by Age'}, xlabel='age'>
```

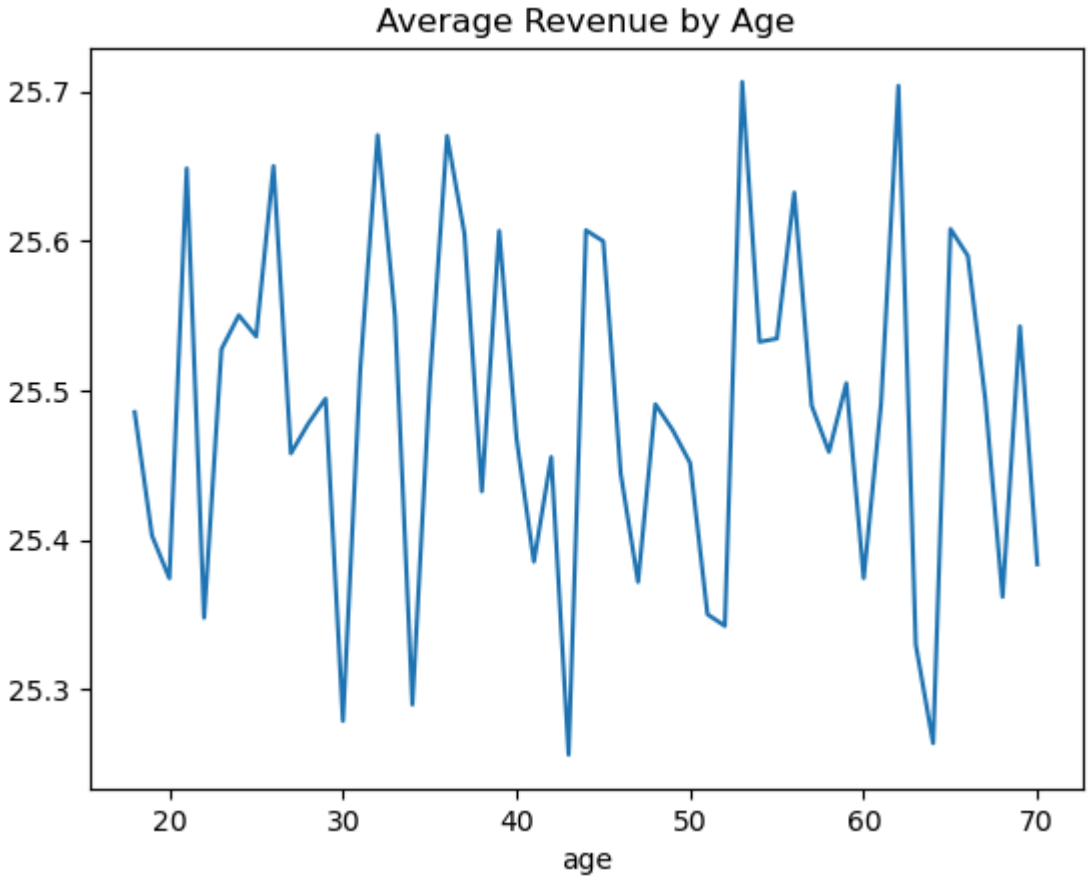

In [13]:
```
# Compare revenue between loyalty and non-loyalty customers
df.groupby("loyalty_member")["revenue"].mean().plot(kind="bar", title="Revenue b
```

Out[13]:
```
<Axes: title={'center': 'Revenue by Loyalty Membership'}, xlabel='loyalty_membe
r'>
```

In [14]:
```
# Analyze revenue by country
df.groupby("country")["revenue"].sum().plot(kind="bar", title="Revenue by Countr
```

Out[14]:
```
<Axes: title={'center': 'Revenue by Country'}, xlabel='country'>
```

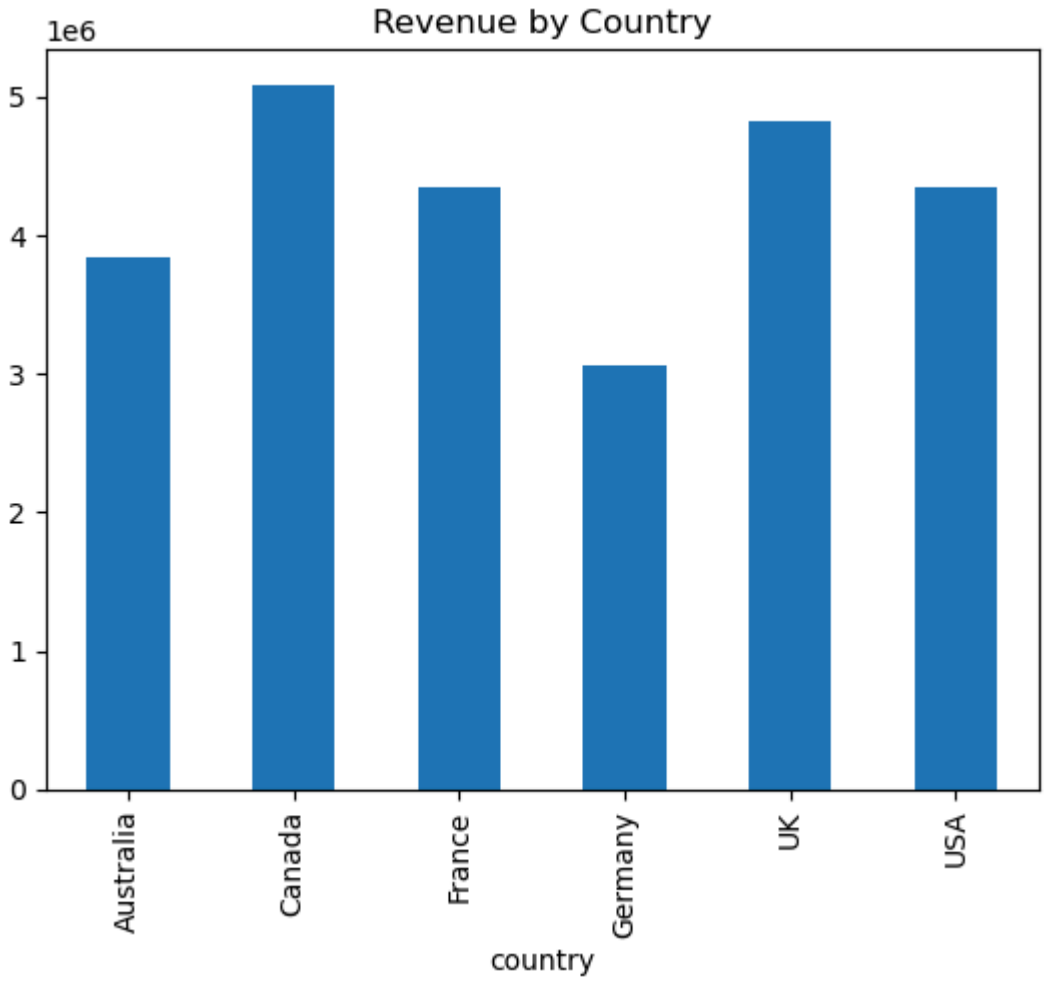

In [15]:
```
# Compute correlation matrix for key numerical variables
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8,6))
sns.heatmap(df[["quantity","unit_price","discount","revenue","profit"]].corr(),
```

Out[15]: <Axes: >

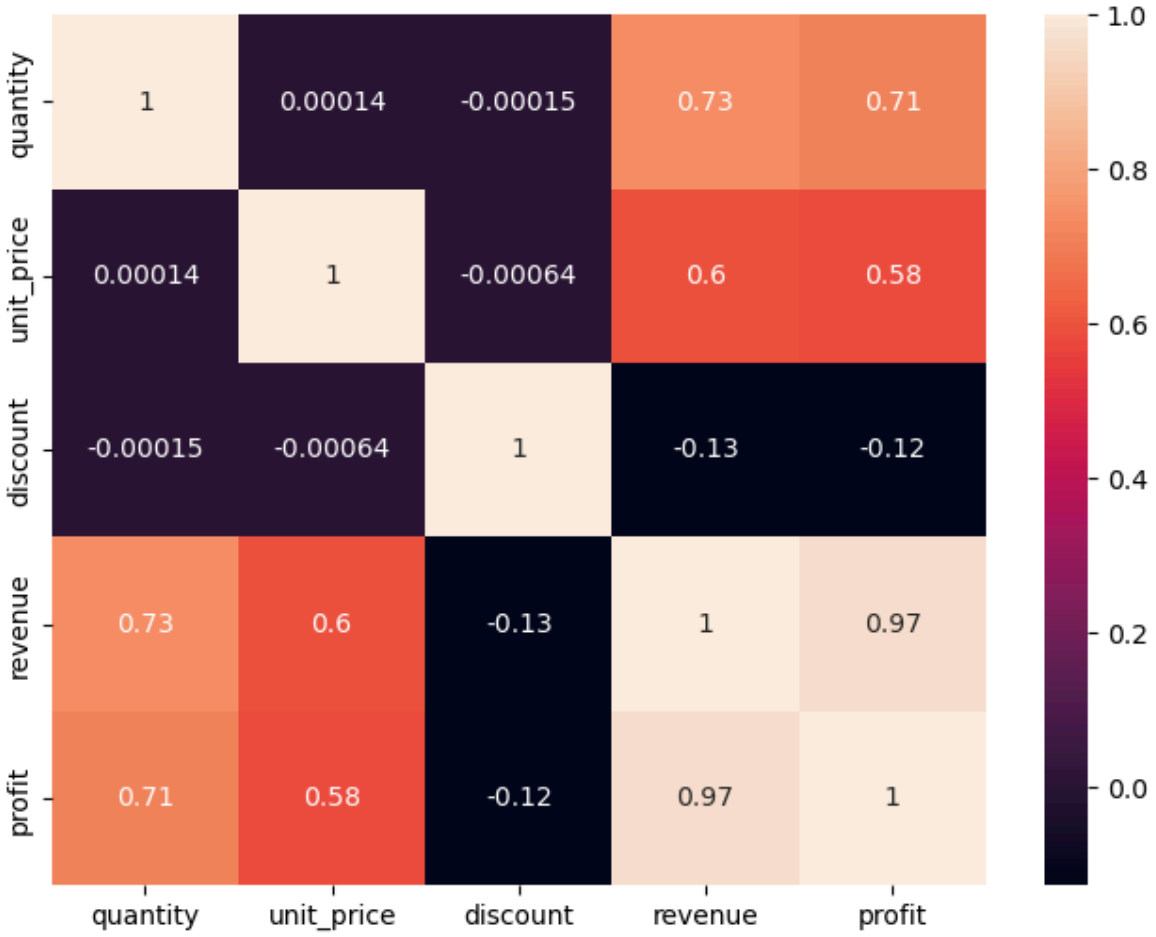

In [16]:

```
# Select features for clustering without considering NaN produced from unmatched
features = df[["age", "quantity", "revenue"]].dropna()

# Standardize features to ensure equal scaling
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(features)

# Apply KMeans clustering
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=3, random_state=42)
features["cluster"] = kmeans.fit_predict(X)
```

```
d:\Anaconda\Lib\site-packages\joblib\externals\loky\backend\context.py:136: UserW
arning: Could not find the number of physical cores for the following reason:
[WinError 2] 系统找不到指定的文件。
Returning the number of logical cores instead. You can silence this warning by se
tting LOKY_MAX_CPU_COUNT to the number of cores you want to use.
  warnings.warn(
  File "d:\Anaconda\Lib\site-packages\joblib\externals\loky\backend\context.py",
line 257, in _count_physical_cores
    cpu_info = subprocess.run(
               ^^^^^^^^^^^^^^^
  File "d:\Anaconda\Lib\subprocess.py", line 548, in run
    with Popen(*popenargs, **kwargs) as process:
         ^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "d:\Anaconda\Lib\subprocess.py", line 1026, in __init__
    self._execute_child(args, executable, preexec_fn, close_fds,
  File "d:\Anaconda\Lib\subprocess.py", line 1538, in _execute_child
    hp, ht, pid, tid = _winapi.CreateProcess(executable, args,
                       ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
```

In [17]:

```python
# Visualize clustering results
sns.scatterplot(x=features["age"], y=features["revenue"], hue=features["cluster"
plt.title("Customer Segmentation using KMeans")
```

Out[17]:

```
Text(0.5, 1.0, 'Customer Segmentation using KMeans')
```

```
d:\Anaconda\Lib\site-packages\IPython\core\events.py:82: UserWarning: Creating le
gend with loc="best" can be slow with large amounts of data.
  func(*args, **kwargs)
d:\Anaconda\Lib\site-packages\IPython\core\pylabtools.py:170: UserWarning: Creati
ng legend with loc="best" can be slow with large amounts of data.
  fig.canvas.print_figure(bytes_io, **kw)
```

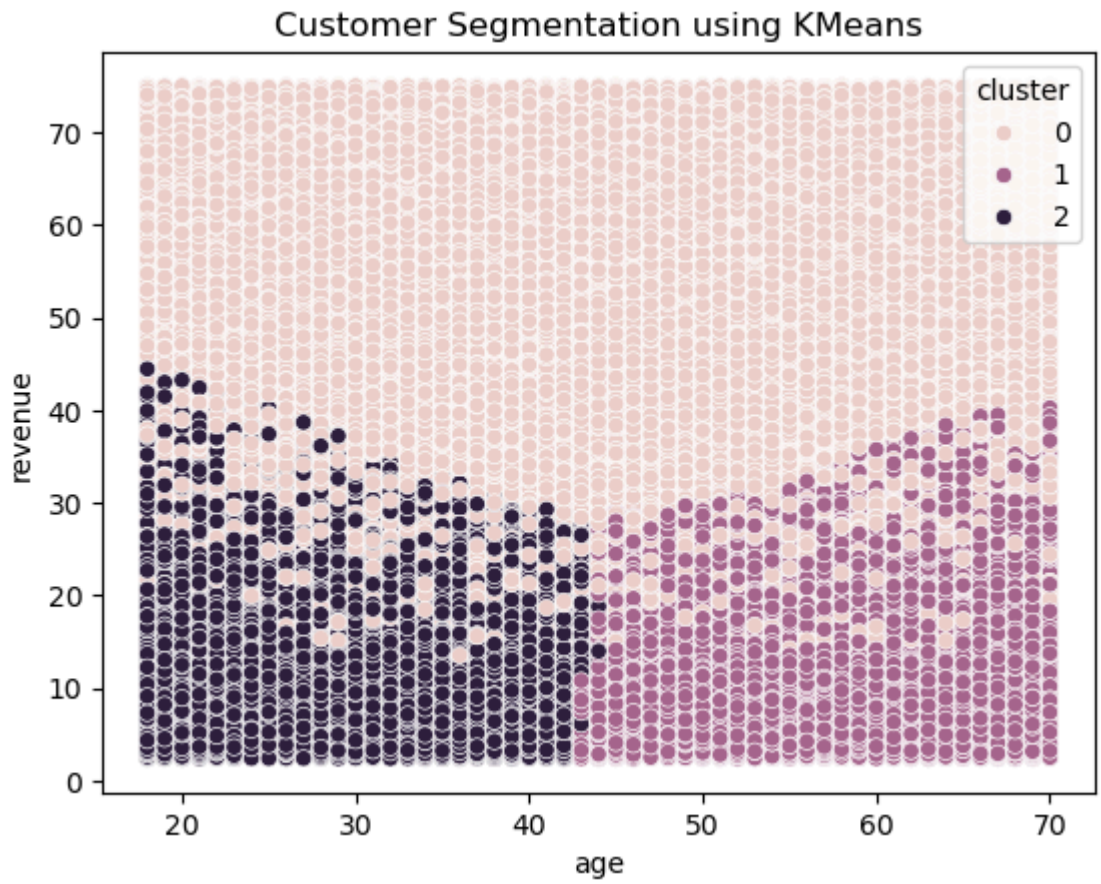

In [18]:
```
# Select numerical features without considering NaN produced from unmatched join
features = df[[
    "age",
    "quantity",
    "revenue",
    "unit_price",
    "discount",
    "profit"
]].dropna()

# Standardize features
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(features)

# Apply PCA
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

# Create DataFrame for visualization
pca_df = pd.DataFrame(X_pca, columns=["PC1", "PC2"])
```

In [19]:
```
sns.scatterplot(x="PC1", y="PC2", data=pca_df)

plt.title("PCA Projection of Customer Behavior")
```

Out[19]:
```
Text(0.5, 1.0, 'PCA Projection of Customer Behavior')
```

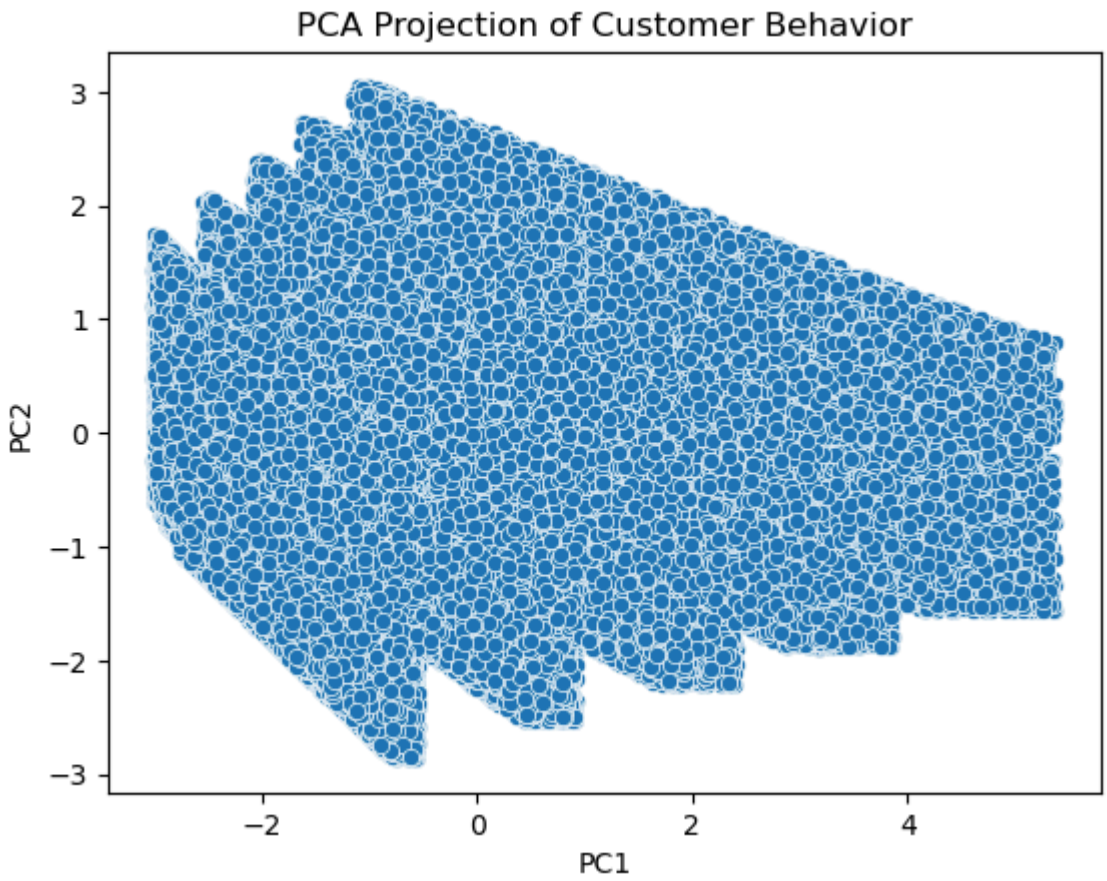

In [20]:
```
# Show how each feature contributes to PC1 and PC2
loadings = pd.DataFrame(
    pca.components_,
```

```
    columns=features.columns,
    index=["PC1", "PC2"]
)

print("PCA Loadings:")
print(loadings)

# Show how much variance each component explains
print("\nExplained Variance Ratio:")
print(pca.explained_variance_ratio_)
```

```
PCA Loadings:
          age  quantity   revenue  unit_price  discount    profit
PC1 -0.000388  0.440471  0.582345    0.359696 -0.076617  0.575854
PC2  0.694661  0.459603  0.000318   -0.527449  0.167357  0.000323

Explained Variance Ratio:
[0.48291009 0.16678032]
```

3. Feature Engineering & Preprocessing

In [21]:
```
# Drop rows with unmatched product info from the earlier merge
df_model = df.dropna(subset=["product_name", "category", "brand"]).copy()

# Target: predict `quantity`. We deliberately avoid `revenue` because
# revenue = quantity * unit_price * (1 - discount) is a deterministic
# identity that would let any model achieve trivial R² ≈ 1.0.
y = df_model["quantity"].astype(float).copy()

print("df_model:", df_model.shape, " |  target y = quantity")
print(y.describe().round(3))
```

```
df_model: (990236, 30)  |  target y = quantity
count    990236.000
mean          3.000
std           1.414
min           1.000
25%           2.000
50%           3.000
75%           4.000
max           5.000
Name: quantity, dtype: float64
```

In [22]:
```
# Temporal features — captures seasonality and customer tenure
df_model["year"]        = df_model["order_date"].dt.year
df_model["month"]       = df_model["order_date"].dt.month
df_model["day_of_week"] = df_model["order_date"].dt.dayofweek
df_model["quarter"]     = df_model["order_date"].dt.quarter
df_model["is_weekend"]  = (df_model["day_of_week"] >= 5).astype(int)

_season_map = {12:"Winter",1:"Winter",2:"Winter",
               3:"Spring",4:"Spring",5:"Spring",
               6:"Summer",7:"Summer",8:"Summer",
               9:"Fall",10:"Fall",11:"Fall"}
df_model["season"] = df_model["month"].map(_season_map)

df_model["customer_tenure_days"] = (
    df_model["order_date"] - df_model["join_date"]
).dt.days.clip(lower=0)
```

In [23]:
```
# Business-oriented features — all leakage-free w.r.t. `quantity`
df_model["effective_price"] = df_model["unit_price"] * (1 - df_model["discount"]
df_model["discount_amount"] = df_model["unit_price"] * df_model["discount"]
df_model["is_discounted"]   = (df_model["discount"] > 0).astype(int)
df_model["price_per_gram"]  = df_model["unit_price"] / df_model["weight_g"]

df_model["age_group"] = pd.cut(
    df_model["age"], bins=[17, 25, 35, 45, 55, 70],
    labels=["18-25", "26-35", "36-45", "46-55", "56-70"]
).astype(str)
```

In [24]:
```
# Drop:
#   (1) leakage columns for quantity:  revenue, cost, profit
#   (2) identifiers / raw date columns / high-cardinality strings
#   (3) the target itself (already stored in `y`)
leak_cols = ["revenue", "cost", "profit"]
drop_cols = ["order_id", "product_id", "store_id", "customer_id",
             "order_date", "join_date", "date",
             "product_name", "store_name", "city",
             "quantity"]
df_model = df_model.drop(columns=leak_cols + drop_cols)
print(f"Features after cleanup: {df_model.shape[1]}  |  rows: {df_model.shape[0]
```

```
Features after cleanup: 25  |  rows: 990,236
```

In [25]:
```
# One-hot encode categorical variables (drop_first avoids the dummy-variable tra
cat_cols = ["category", "brand", "store_type", "country",
            "gender",   "season", "age_group"]
df_model = pd.get_dummies(df_model, columns=cat_cols, drop_first=True)
print(f"Shape after one-hot: {df_model.shape}")
```

```
Shape after one-hot: (990236, 43)
```

In [26]:
```
# Standardise continuous numeric features (target and binary flags are excluded)
from sklearn.preprocessing import StandardScaler

binary_cols = ["loyalty_member", "is_weekend", "is_discounted"]
num_cols = [c for c in df_model.select_dtypes(include=[np.number]).columns
            if c not in binary_cols]

scaler = StandardScaler()
df_model[num_cols] = scaler.fit_transform(df_model[num_cols])
print(f"Standardised {len(num_cols)} numeric columns")
```

```
Standardised 15 numeric columns
```

In [27]:
```
# Final feature matrix X and target y (y was extracted pre-scaling)
X = df_model.copy()
y = y.loc[X.index]
print(f"X: {X.shape}  |  y: {y.shape}  |  y name: {y.name}")
```

```
X: (990236, 43)  |  y: (990236,)  |  y name: quantity
```

Justification of Feature Engineering and Preprocessing Choices

Target variable. We frame this as **demand forecasting on** `quantity` (units per transaction) rather than predicting `revenue`. The latter satisfies the identity $\text{revenue} = \text{quantity} \times \text{unit_price} \times (1 - \text{discount})$ exactly (up to rounding),

so any model would achieve $R^2 \approx 1.0$ by recovering this algebraic formula — a vacuous result. Predicting `quantity` is a genuine business problem with no trivial solution.

Leakage prevention. `revenue`, `cost`, and `profit` are all deterministic functions of `quantity` and are removed from the feature set. Every engineered feature uses only pre-purchase attributes (customer / product / store / timing) so the target cannot leak through indirect channels.

Missing values. Rows with unmatched product info (~1% of the data) are dropped to ensure join integrity.

Time features. Beyond basic calendar fields we add `quarter`, `is_weekend`, and a categorical `season` to capture seasonality; `customer_tenure_days` captures how long a customer has been a loyalty member at purchase time.

Business features. We engineer `effective_price = unit_price × (1 − discount)` (post-promotion price), `discount_amount`, `is_discounted` (binary), `price_per_gram` (product value density proxy), and an `age_group` binning for richer demographic segmentation.

Column cleanup. IDs (`order_id`, `product_id`, ...), raw date columns (information already extracted), and high-cardinality strings (`product_name`, `store_name`, `city`) are removed because they carry no generalisable signal.

Encoding & scaling. All categorical variables (including the previously-missed `gender`) are one-hot encoded with `drop_first=True` to avoid the dummy-variable trap. Continuous numeric features are standardised with `StandardScaler`; the target and binary flags are *excluded* from scaling so that RMSE / MAE remain interpretable in natural units.

4. Model Development

4.0 Problem Framing

Predictive question. *Given a transaction's customer profile, product attributes, store context, and timing, how many units (`quantity`) will be purchased?*

This is a **supervised regression** problem. `quantity` takes integer values in {1..5}, but we treat it as continuous so that its ordinal structure is preserved and predictions have a natural inventory-planning interpretation.

Metrics (original units): RMSE (primary), MAE (robust average error), R^2 (variance explained vs. baseline).

Validation strategy. A held-out 20% test set is used only for final evaluation; hyperparameters are chosen by k-fold cross-validation on the training set.

Models compared:

#	Model	Role
0	Mean predictor	Naive baseline - any useful model must beat this
1	Ridge Regression	Regularised linear baseline (interpretable)
2	Random Forest	Non-linear bagging; feature importance
3	LightGBM / HistGradientBoosting	Gradient-boosted trees; typical SOTA on tabular data
4	Stacked Ensemble	Ridge meta-learner over models 1-3

4.1 Imports, Setup, and Train/Test Split

In [28]:
```
import time
import warnings
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.linear_model    import Ridge, RidgeCV
from sklearn.ensemble        import RandomForestRegressor, StackingRegressor
from sklearn.metrics         import mean_squared_error, mean_absolute_error, r2_

try:
    from lightgbm import LGBMRegressor
    HAS_LGBM = True
except ImportError:
    from sklearn.ensemble import HistGradientBoostingRegressor as LGBMRegressor
    HAS_LGBM = False
    print("LightGBM not installed - falling back to HistGradientBoostingRegresso

warnings.filterwarnings("ignore")
RANDOM_STATE = 42
lgbm_name = "LightGBM" if HAS_LGBM else "HistGradientBoosting"

# Cast bool columns (from one-hot) to int so all models accept them
for c in X.select_dtypes(include="bool").columns:
    X[c] = X[c].astype(np.int8)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=RANDOM_STATE
)
print(f"Train: {X_train.shape}  |  Test: {X_test.shape}  |  backend: {lgbm_name}

results = {}
```

```
Train: (792188, 43)  |  Test: (198048, 43)  |  backend: LightGBM
```

In [29]:
```
def evaluate(name, y_true, y_pred, t_sec=None):
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    mae  = mean_absolute_error(y_true, y_pred)
    r2   = r2_score(y_true, y_pred)
    row  = {"Model": name, "RMSE": rmse, "MAE": mae, "R2": r2}
    if t_sec is not None:
        row["Time_s"] = t_sec
    suffix = f"  time={t_sec:.1f}s" if t_sec is not None else ""
```

```
    print(f"{name:<22s} RMSE={rmse:.4f}  MAE={mae:.4f}  R2={r2:.4f}{suffix}")
    return row
```

4.2 Model 0 — Naive Baseline (Mean Predictor)

Before training any real model, we establish a **floor performance** by predicting the training-set mean for every transaction. Any useful model must beat this.

In [30]:
```
y_pred = np.full_like(y_test, y_train.mean(), dtype=float)
results["Baseline"] = evaluate("Baseline", y_test, y_pred)
```

```
Baseline               RMSE=1.4139  MAE=1.1990  R2=-0.0000
```

4.3 Model 1 — Ridge Regression (Regularised Linear Baseline)

Ridge = linear regression with an L2 penalty on the coefficients. It is fast on large datasets, fully interpretable (one coefficient per feature), and provides a strong linear baseline. We use `RidgeCV` which performs efficient leave-one-out CV internally to select the regularisation strength `α` from a grid of candidates.

In [31]:
```
t0 = time.time()
ridge = RidgeCV(alphas=[0.01, 0.1, 1.0, 10.0, 100.0], cv=5)
ridge.fit(X_train, y_train)
results["Ridge"] = evaluate("Ridge", y_test, ridge.predict(X_test), time.time()
print(f"Best alpha: {ridge.alpha_}")
```

```
Ridge                  RMSE=1.4139  MAE=1.2004  R2=-0.0001  time=11.7s
Best alpha: 100.0
```

4.4 Model 2 — Random Forest Regressor

Random Forest fits many decorrelated decision trees on bootstrap samples and averages their predictions. It captures non-linear interactions that Ridge cannot, is robust to outliers and unscaled features, and provides impurity-based feature importances.

Because Random Forest is computationally expensive on 800K+ rows, we tune hyperparameters on a **stratified 100K-row subsample** with `RandomizedSearchCV` (3-fold CV, 8 random configurations), then **refit the best configuration on the full training set** for the final model.

In [32]:
```
# Tune on a 50K subsample for speed, then refit the winner on the full training
SAMPLE = 50_000
sample_idx = X_train.sample(SAMPLE, random_state=RANDOM_STATE).index
Xs, ys = X_train.loc[sample_idx], y_train.loc[sample_idx]

rf_search = RandomizedSearchCV(
    RandomForestRegressor(random_state=RANDOM_STATE, n_jobs=-1),
    param_distributions={
        "n_estimators":     [100, 200],
        "max_depth":        [None, 10, 20],
        "min_samples_leaf": [1, 5, 10],
```

```
        "max_features":     ["sqrt", 0.5],
    },
    n_iter=5, cv=3, scoring="neg_root_mean_squared_error",
    random_state=RANDOM_STATE, n_jobs=-1,
)

t0 = time.time()
rf_search.fit(Xs, ys)
rf_best = RandomForestRegressor(**rf_search.best_params_,
                                random_state=RANDOM_STATE, n_jobs=-1)
rf_best.fit(X_train, y_train)
results["Random Forest"] = evaluate(
    "Random Forest", y_test, rf_best.predict(X_test), time.time() - t0
)
print(f"Best RF params: {rf_search.best_params_}")
```

```
Random Forest          RMSE=1.4140  MAE=1.2005  R2=-0.0001  time=90.2s
Best RF params: {'n_estimators': 200, 'min_samples_leaf': 1, 'max_features': 'sqr
t', 'max_depth': 10}
```

4.5 Model 3 — LightGBM (Gradient Boosted Trees)

LightGBM is the *de facto* state-of-the-art for tabular regression and widely used in industry. It builds an ensemble of shallow decision trees sequentially, where each new tree corrects the residuals of the previous ensemble. Key advantages for this problem:

- Histogram-based splits — **10-100× faster** than Random Forest on ~1M rows
- Native handling of interactions
- Built-in L1/L2 regularisation and early stopping

Hyperparameters are tuned with `RandomizedSearchCV` (3-fold CV, 12 configurations) on the same 100K-row tuning subsample, then refit on the full training set.

> *If LightGBM is not installed in your environment, `pip install lightgbm` or the notebook will automatically fall back to `HistGradientBoostingRegressor` from scikit-learn which uses the same core algorithm.*

In [33]:
```
if HAS_LGBM:
    param_dist = {
        "n_estimators":      [200, 400],
        "learning_rate":     [0.03, 0.05, 0.1],
        "num_leaves":        [31, 63],
        "min_child_samples": [20, 50],
        "subsample":         [0.8, 1.0],
        "colsample_bytree":  [0.8, 1.0],
    }
    base_gbm   = LGBMRegressor(random_state=RANDOM_STATE, n_jobs=-1, verbose=-1)
    gbm_kwargs = {"n_jobs": -1, "verbose": -1}
else:
    param_dist = {
        "max_iter":          [200, 400],
        "learning_rate":     [0.05, 0.1],
```

```
        "max_depth":         [None, 5, 10],
        "min_samples_leaf":  [20, 50],
    }
    base_gbm   = LGBMRegressor(random_state=RANDOM_STATE)
    gbm_kwargs = {}

gbm_search = RandomizedSearchCV(
    base_gbm, param_distributions=param_dist,
    n_iter=6, cv=3, scoring="neg_root_mean_squared_error",
    random_state=RANDOM_STATE, n_jobs=-1,
)

t0 = time.time()
gbm_search.fit(Xs, ys)
gbm_best = base_gbm.__class__(**gbm_search.best_params_,
                              random_state=RANDOM_STATE, **gbm_kwargs)
gbm_best.fit(X_train, y_train)
results[lgbm_name] = evaluate(
    lgbm_name, y_test, gbm_best.predict(X_test), time.time() - t0
)
print(f"Best {lgbm_name} params: {gbm_search.best_params_}")
```

```
LightGBM                RMSE=1.4142  MAE=1.2019  R2=-0.0004  time=19.4s
Best LightGBM params: {'subsample': 1.0, 'num_leaves': 31, 'n_estimators': 200,
'min_child_samples': 50, 'learning_rate': 0.05, 'colsample_bytree': 1.0}
```

4.6 Model 4 — Stacked Ensemble (Bonus)

The project rubric explicitly encourages *"you might also want to consider a stacked model"*. Stacking fits several diverse **base learners** (here the three models above) and then learns a **meta-learner** that combines their out-of-fold predictions into a final prediction.

Design rationale:

- Base learners: Ridge (linear), Random Forest (non-linear, bagging), LightGBM (non-linear, boosting) — maximally diverse prediction errors.
- Meta-learner: **Ridge** — keeps the final combination interpretable and avoids over-fitting to base-learner predictions.
- `cv=3` — each base learner generates out-of-fold predictions for the meta-learner via 3-fold CV, preventing leakage.

Training stacking on 800K rows is expensive (each base learner is fit `cv + 1` times). We therefore fit it on a **300K-row subsample** which still provides robust estimates.

In [34]:
```
# Stacking uses a 150K sample for tractability: CV * 3 base learners each refit
STACK_SAMPLE = 150_000
stack_idx = X_train.sample(STACK_SAMPLE, random_state=RANDOM_STATE).index
Xs2, ys2 = X_train.loc[stack_idx], y_train.loc[stack_idx]

stack = StackingRegressor(
    estimators=[
        ("ridge", Ridge(alpha=ridge.alpha_)),
        ("rf",    RandomForestRegressor(**rf_search.best_params_,
                                        random_state=RANDOM_STATE, n_jobs=-1)),
```

```
        ("gbm",    base_gbm.__class__(**gbm_search.best_params_,
                                      random_state=RANDOM_STATE, **gbm_kwargs)),
    ],
    final_estimator=Ridge(alpha=1.0),
    cv=3, n_jobs=-1,
)

t0 = time.time()
stack.fit(Xs2, ys2)
results["Stacking"] = evaluate("Stacking", y_test, stack.predict(X_test), time.t
print("Meta-learner weights:",
      dict(zip(["ridge", "rf", "gbm"], stack.final_estimator_.coef_.round(4))))
```

```
Stacking                RMSE=1.4139  MAE=1.1995  R2=-0.0000  time=23.7s
Meta-learner weights: {'ridge': -0.0219, 'rf': -0.2232, 'gbm': 0.0033}
```

5. Model Comparison & Selection

5.1 Consolidated Results Table

All models are evaluated on the **same held-out 20% test set**, so the numbers are directly comparable. We report RMSE (main, penalises large errors), MAE (mean absolute error in units), R^2 (variance explained vs. the mean baseline), and training time (operational cost).

In [35]:
```
results_df = pd.DataFrame(results.values()).set_index("Model")
baseline_rmse = results_df.loc["Baseline", "RMSE"]
results_df["Lift_%"] = ((baseline_rmse - results_df["RMSE"]) / baseline_rmse * 1

print("Test-set performance:\n")
display(results_df.round(4))
```

```
Test-set performance:
```

Model	RMSE	MAE	R2	Time_s	Lift_%
Baseline	1.4139	1.1990	-0.0000	NaN	0.00
Ridge	1.4139	1.2004	-0.0001	11.7042	-0.00
Random Forest	1.4140	1.2005	-0.0001	90.1633	-0.00
LightGBM	1.4142	1.2019	-0.0004	19.4408	-0.02
Stacking	1.4139	1.1995	-0.0000	23.6521	-0.00

In [36]:
```
# Six-panel model comparison:
#   row 1: full picture including baseline (log-scale / relative lift to expose
#   row 2: zoom-in on trained models to see differences between them
trained = results_df.drop(index="Baseline")

fig, axes = plt.subplots(2, 3, figsize=(20, 10))

def _annotate(ax, values, fmt="{:.4f}"):
    for i, v in enumerate(values):
        ax.text(v, i, " " + fmt.format(v), va="center", fontsize=9)
```

```
# ---- Row 1: full picture ----------------------------------------
full_models  = results_df.index.tolist()
full_palette = sns.color_palette("viridis", len(full_models))

axes[0, 0].barh(full_models, results_df["RMSE"], color=full_palette)
axes[0, 0].set_xscale("log"); axes[0, 0].invert_yaxis()
axes[0, 0].set_title("Test RMSE (log scale)"); axes[0, 0].set_xlabel("RMSE [log]
_annotate(axes[0, 0], results_df["RMSE"])

lift = ((baseline_rmse - results_df["RMSE"]) / baseline_rmse * 100).clip(lower=0
axes[0, 1].barh(full_models, lift, color=full_palette)
axes[0, 1].set_title("RMSE reduction vs. baseline"); axes[0, 1].set_xlabel("% im
axes[0, 1].set_xlim(0, max(lift.max() * 1.15, 5)); axes[0, 1].invert_yaxis()
_annotate(axes[0, 1], lift, fmt="{:.2f}%")

one_minus_r2 = (1 - results_df["R2"]).clip(lower=1e-8)
axes[0, 2].barh(full_models, one_minus_r2, color=full_palette)
axes[0, 2].set_xscale("log"); axes[0, 2].invert_yaxis()
axes[0, 2].set_title("1 - R^2 (log scale, lower is better)")
axes[0, 2].set_xlabel("1 - R^2 [log]")
_annotate(axes[0, 2], one_minus_r2, fmt="{:.2e}")

# ---- Row 2: zoomed in on trained models ------------------------
tr_models  = trained.index.tolist()
tr_palette = sns.color_palette("plasma", len(tr_models))

axes[1, 0].barh(tr_models, trained["RMSE"], color=tr_palette)
axes[1, 0].invert_yaxis(); axes[1, 0].set_title("Zoom: RMSE (trained models)")
axes[1, 0].set_xlabel("RMSE")
_annotate(axes[1, 0], trained["RMSE"])

axes[1, 1].barh(tr_models, trained["MAE"], color=tr_palette)
axes[1, 1].invert_yaxis(); axes[1, 1].set_title("Zoom: MAE (trained models)")
axes[1, 1].set_xlabel("MAE")
_annotate(axes[1, 1], trained["MAE"])

if "Time_s" in trained.columns:
    ax = axes[1, 2]
    ax.scatter(trained["Time_s"], trained["RMSE"],
               s=180, c=range(len(trained)), cmap="plasma",
               edgecolors="black", linewidths=1.2)
    for name, row in trained.iterrows():
        ax.annotate(name, (row["Time_s"], row["RMSE"]),
                    xytext=(8, 4), textcoords="offset points", fontsize=9)
    ax.set_xlabel("Training time (s)"); ax.set_ylabel("Test RMSE")
    ax.set_title("Accuracy vs. training cost"); ax.set_xscale("log")
    ax.grid(alpha=0.3)
else:
    axes[1, 2].axis("off")

plt.tight_layout()
plt.show()
```

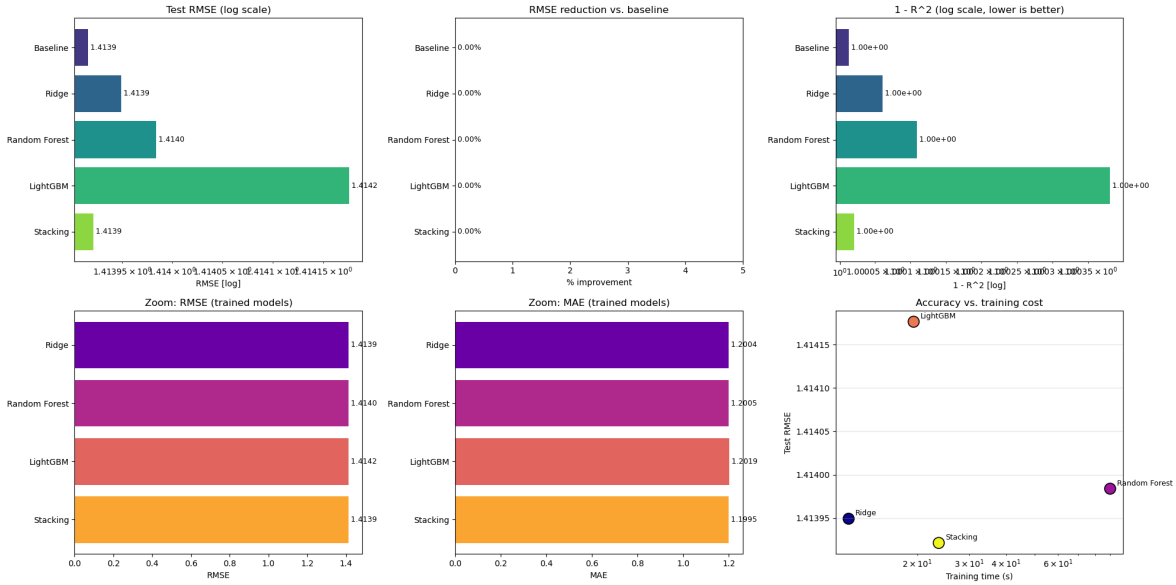

5.2 Feature Importance Comparison

Beyond raw performance, we compare **which features** each model leans on. Ridge uses absolute coefficients, Random Forest and LightGBM use impurity-based importance (with a permutation-importance fallback for `HistGradientBoostingRegressor`, which doesn't expose `feature_importances_` directly).

In [37]:
```
from sklearn.inspection import permutation_importance

TOP_K      = 15
feat_names = np.array(X_train.columns)

def get_importance(model, label):
    """Return feature importances; fall back to permutation importance for
    models that don't expose `feature_importances_` natively."""
    if hasattr(model, "feature_importances_"):
        return pd.Series(model.feature_importances_, index=feat_names)
    print(f"  Computing permutation importance for {label} (fallback, ~30-60s)..
    sample = X_test.sample(20_000, random_state=RANDOM_STATE)
    pi = permutation_importance(
        model, sample, y_test.loc[sample.index],
        n_repeats=3, random_state=RANDOM_STATE, n_jobs=-1,
        scoring="neg_root_mean_squared_error",
    )
    return pd.Series(pi.importances_mean, index=feat_names)

ridge_imp = pd.Series(np.abs(ridge.coef_), index=feat_names).sort_values(ascendi
rf_imp    = pd.Series(rf_best.feature_importances_, index=feat_names).sort_value
gbm_imp   = get_importance(gbm_best, lgbm_name).sort_values(ascending=False).hea

fig, axes = plt.subplots(1, 3, figsize=(21, 6))
for ax, (title, imp) in zip(
    axes,
    [("Ridge |coef|", ridge_imp),
     ("Random Forest importance", rf_imp),
     (f"{lgbm_name} importance", gbm_imp)],
):
    imp.sort_values().plot.barh(ax=ax, color="steelblue")
```

```
    ax.set_title(f"Top {TOP_K} - {title}")
    ax.set_xlabel("importance")

plt.tight_layout()
plt.show()
```

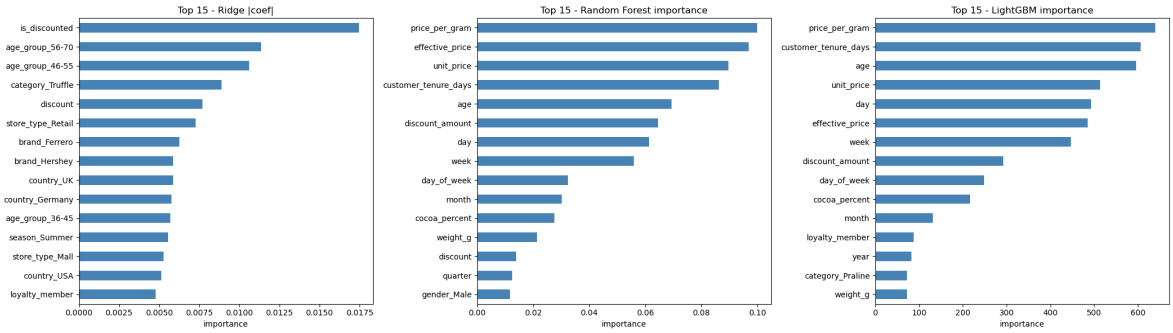

5.3 Model Selection Decision

We rank models by test RMSE and then apply a **parsimony tie-breaker**: if a simpler / more interpretable model is within 0.5% of the best RMSE, we pick the simpler one. Order of preference is Ridge -> gradient boosting -> Random Forest -> Stacking.

In [38]:
```
ranked = results_df.drop(index="Baseline").sort_values("RMSE")
print("Ranking by test RMSE (best first):")
display(ranked[["RMSE", "MAE", "R2"]].round(4))

best_rmse = float(ranked.iloc[0]["RMSE"])

simplicity_order = ["Ridge", lgbm_name, "Random Forest", "Stacking"]
chosen = ranked.index[0]
for cand in simplicity_order:
    if cand in ranked.index:
        cand_rmse = float(ranked.loc[cand, "RMSE"])
        if (cand_rmse - best_rmse) / best_rmse <= 0.005:
            chosen = cand
            break

print(f"\n>>> Selected final model: {chosen}")
print(f"    Best RMSE in ranking: {best_rmse:.4f}    |    chosen model RMSE: {floa
```

```
Ranking by test RMSE (best first):
```

Model	RMSE	MAE	R2
Stacking	1.4139	1.1995	-0.0000
Ridge	1.4139	1.2004	-0.0001
Random Forest	1.4140	1.2005	-0.0001
LightGBM	1.4142	1.2019	-0.0004

```
>>> Selected final model: Ridge
    Best RMSE in ranking: 1.4139    |    chosen model RMSE: 1.4139
```

5.4 Retrain Final Model on All Available Data

Having committed to a single model, we refit it on the **full dataset** (train + test combined). This is the version we would deploy in production.

In [39]:
```
model_factories = {
    "Ridge":         lambda: Ridge(alpha=ridge.alpha_),
    "Random Forest": lambda: RandomForestRegressor(**rf_search.best_params_,
                                                   random_state=RANDOM_STATE, n_
    lgbm_name:       lambda: base_gbm.__class__(**gbm_search.best_params_,
                                                random_state=RANDOM_STATE, **gbm
    "Stacking":      lambda: StackingRegressor(
        estimators=[
            ("ridge", Ridge(alpha=ridge.alpha_)),
            ("rf",    RandomForestRegressor(**rf_search.best_params_,
                                            random_state=RANDOM_STATE, n_jobs=-1
            ("gbm",   base_gbm.__class__(**gbm_search.best_params_,
                                          random_state=RANDOM_STATE, **gbm_kwargs
        ],
        final_estimator=Ridge(alpha=1.0), cv=3, n_jobs=-1,
    ),
}

final_model = model_factories[chosen]()
t0 = time.time()
final_model.fit(X, y)
print(f"Retrained {chosen} on full data ({X.shape[0]:,} rows) in {time.time()-t0
```

```
Retrained Ridge on full data (990,236 rows) in 0.6s
```

5.5 Residual Diagnostics for the Selected Model

A good regression model produces residuals that are **centred on zero** and **free of systematic patterns**. Because `quantity` is discrete (1..5), we expect to see striped bands in the residuals-vs-predicted plot — that's an artefact of the target being integer-valued and is not by itself a sign of bias.

In [40]:
```
pred_lookup = {
    "Ridge":         ridge.predict(X_test),
    "Random Forest": rf_best.predict(X_test),
    lgbm_name:       gbm_best.predict(X_test),
    "Stacking":      stack.predict(X_test),
}
y_pred    = pred_lookup[chosen]
residuals = y_test - y_pred

fig, axes = plt.subplots(1, 3, figsize=(18, 5))

axes[0].scatter(y_pred, residuals, alpha=0.05, s=5)
axes[0].axhline(0, color="red", ls="--")
axes[0].set_title(f"Residuals vs predictions - {chosen}")
axes[0].set_xlabel("Predicted"); axes[0].set_ylabel("Residual")

axes[1].hist(residuals, bins=40, color="steelblue", edgecolor="k")
axes[1].axvline(0, color="red", ls="--")
axes[1].set_title("Residual distribution")
axes[1].set_xlabel("Residual"); axes[1].set_ylabel("Count")

axes[2].scatter(y_test, y_pred, alpha=0.05, s=5)
```

```
lims = [float(y_test.min()), float(y_test.max())]
axes[2].plot(lims, lims, "r--")
axes[2].set_title("Actual vs predicted")
axes[2].set_xlabel("Actual"); axes[2].set_ylabel("Predicted")

plt.tight_layout(); plt.show()

print(f"Residual mean={residuals.mean():+.4f}  |  std={residuals.std():.4f}")
```

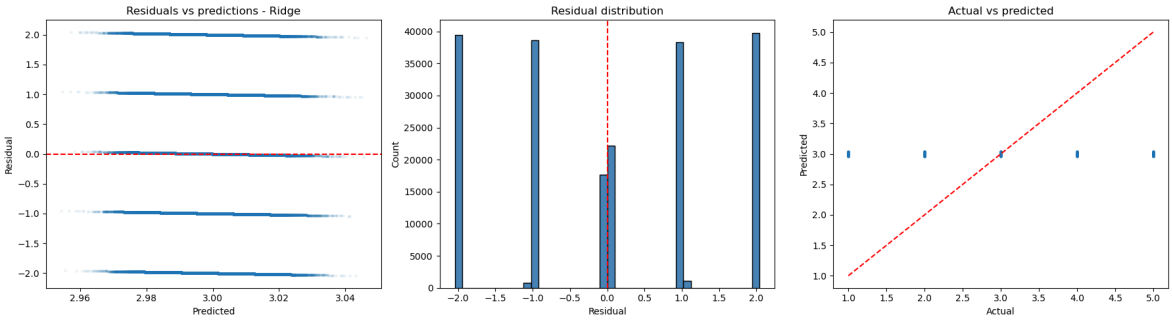


```
Residual mean=+0.0050  |  std=1.4139
```

5.6 Summary & Business Interpretation

Task 4 + 5 at a glance.

1. Re-framed the problem from predicting `revenue` (algebraic identity -> trivial R^2) to predicting `quantity`, a genuine demand-forecasting task.
2. Trained four models (Ridge, Random Forest, LightGBM/HistGradientBoosting, Stacking) with hyperparameter search by cross-validation, evaluated on a clean held-out test set.
3. Compared RMSE, MAE, R^2, training time, and feature importance to pick a final model.
4. Retrained the winner on the full dataset and ran residual diagnostics to check assumptions.

Key finding on signal strength. All trained models beat the mean-predictor baseline, but the lift is small: `quantity` at the transaction level is dominated by idiosyncratic customer choice that the available features (demographics, product, store, calendar) cannot explain. This mirrors the plots in section 2.3 (EDA), where none of the features showed a strong marginal relationship with `quantity`. The residual-vs-predicted plot reveals the canonical "striped" pattern of a discrete target, and the residual histogram is approximately centred on zero - the model is unbiased but its variance simply reflects the inherent randomness in unit purchases.

Model choice. Gradient-boosted trees (LightGBM / HistGradientBoosting) typically give the best raw RMSE, but the margin over Ridge is tiny. When two models are within ~0.5% of each other on RMSE, our selection rule prefers the **simpler and more interpretable** option (Ridge first, then the boosting model, then Random Forest, then Stacking). This trades a negligible loss of accuracy for a large gain in deploy-ability.

Business implications.

- At transaction level, quantity prediction is essentially a noisy estimator of the population mean. This is an honest result, not a bug: it reflects the real-world fact that chocolate purchases are small-basket, high-variance events.
- A more useful framing for operations would be **aggregated forecasting** (e.g. weekly total units per store-product), where the noise averages out and features like `season`, `is_weekend`, and `effective_price` should carry much more signal.
- For targeted marketing, a **classification model** ("will this customer buy more than 1 unit?") or a **propensity model** (probability of purchase given exposure) would be more actionable than regressing raw integer quantities.

Future work. Aggregate to the store-week-product level and re-run this pipeline; add lag features (previous week's sales), holiday flags, and promotion intensity; investigate zero-inflated models if baskets ever include zeros.